



作业习题评讲课

数据库系统原理课程

2024

第一次作业 2.1

2.1 Consider the employee database of Figure 2.17. What are the appropriate primary keys?

employee (*person_name*, *street*, *city*)
works (*person_name*, *company_name*, *salary*)
company (*company_name*, *city*)

Figure 2.17 Employee database.

employee(*person_name*, *street*, *city*)
works(*person_name*, *company_name*, *salary*)
company(*company_name*, *city*)

第一次作业 2.6

employee (*person_name*, *street*, *city*)
works (*person_name*, *company_name*, *salary*)
company (*company_name*, *city*)

Figure 2.17 Employee database.

2.6 Consider the employee database of Figure 2.17. Give an expression in the relational algebra to express each of the following queries:

- Find the name of each employee who lives in city “Miami”.
- Find the name of each employee whose salary is greater than \$100000.
- Find the name of each employee who lives in “Miami” and whose salary is greater than \$100000.

答案:

- $\Pi_{person_name}(\sigma_{city="Miami"}(employee))$
- $\Pi_{person_name}(\sigma_{salary>100000}(works))$
- $\Pi_{person_name}(\sigma_{salary>100000 \wedge city="Miami"}(employee \bowtie works))$

第一次作业 2.7

2.7 Consider the bank database of Figure 2.18. Give an expression in the relational algebra for each of the following queries:

- a. Find the name of each branch located in “Chicago”.
- b. Find the ID of each borrower who has a loan in branch “Downtown”.

branch(*branch_name*, *branch_city*, *assets*)
customer (*ID*, *customer_name*, *customer_street*, *customer_city*)
loan (*loan_number*, *branch_name*, *amount*)
borrower (*ID*, *loan_number*)
account (*account_number*, *branch_name*, *balance*)
depositor (*ID*, *account_number*)

Figure 2.18 Bank database.

答案:

- a. $\Pi_{branch_name}(\sigma_{branch_city="Chicago"}(branch))$
- b. $\Pi_{ID}(\sigma_{branch_city="Downtown"}(borrower \bowtie loan))$

第一次作业 2.8

employee (*person_name*, *street*, *city*)
works (*person_name*, *company_name*, *salary*)
company (*company_name*, *city*)

Figure 2.17 Employee database.

- 2.8 Consider the employee database of Figure 2.17. Give an expression in the relational algebra to express each of the following queries:
- Find the ID and name of each employee who does not work for “BigBank”.
 - Find the ID and name of each employee who earns at least as much as every employee in the database.

答案:

- a. $\Pi_{ID, person_name}(works) - \Pi_{ID, person_name}(\sigma_{company_name = "BigBank"}(works))$
- b. $\Pi_{ID, person_name}(works) - \Pi_{ID, person_name}(\sigma_{works.salary < d.salary}(works \times \rho_d(\Pi_{salary}(works))))$

第一次作业 2.15

branch(*branch_name*, *branch_city*, *assets*)
customer (*ID*, *customer_name*, *customer_street*, *customer_city*)
loan (*loan_number*, *branch_name*, *amount*)
borrower (*ID*, *loan_number*)
account (*account_number*, *branch_name*, *balance*)
depositor (*ID*, *account_number*)

Figure 2.18 Bank database.

- 2.15** Consider the bank database of Figure 2.18. Give an expression in the relational algebra for each of the following queries:
- Find each loan number with a loan amount greater than \$10000.
 - Find the ID of each depositor who has an account with a balance greater than \$6000.
 - Find the ID of each depositor who has an account with a balance greater than \$6000 at the “Uptown” branch.

答案:

- $\Pi_{loan_number}(\sigma_{amount > 10000}(loan))$
- $\Pi_{ID}(\sigma_{balance > 6000}(depositor \bowtie account))$
- $\Pi_{ID}(\sigma_{balance > 6000 \wedge branch_name = "Uptown"}(depositor \bowtie account))$

第二次作业 3.8

- 3.8 Consider the bank database of Figure 3.18, where the primary keys are underlined. Construct the following SQL queries for this relational database.

branch(*branch_name*, *branch_city*, *assets*)
customer (*ID*, *customer_name*, *customer_street*, *customer_city*)
loan (*loan_number*, *branch_name*, *amount*)
borrower (*ID*, *loan_number*)
account (*account_number*, *branch_name*, *balance*)
depositor (*ID*, *account_number*)

Figure 3.18 Banking database.

- a. Find the ID of each customer of the bank who has an account but not a loan.

(select ID
from depositor)
except
(select ID
from borrower)

第二次作业 3.8

b. Find the ID of each customer who lives on the same street and in the same city as customer '12345'.

```
select F.ID  
from customer as F, customer as S  
where F.customer_street = S.customer_street  
and F.customer_city = S.customer_city  
and S.customer_id = '12345'
```

c. Find the name of each branch that has at least one customer who has an account in the bank and who lives in "Harrison".

```
select distinct branch_name  
from account, depositor, customer  
where customer.id = depositor.id  
and depositor.account_number = account.account_number  
and customer_city = 'Harrison'
```


第二次作业 3.9

3.9 Consider the relational database of Figure 3.19, where the primary keys are underlined. Give an expression in SQL for each of the following queries.

- Find the ID, name, and city of residence of each employee who works for “First Bank Corporation”.
- Find the ID, name, and city of residence of each employee who works for “First Bank Corporation” and earns more than \$10000.
- Find the ID of each employee who does not work for “First Bank Corporation”.
- Find the ID of each employee who earns more than every employee of “Small Bank Corporation”.
- Assume that companies may be located in several cities. Find the name of each company that is located in every city in which “Small Bank Corporation” is located.
- Find the name of the company that has the most employees (or companies, in the case where there is a tie for the most).
- Find the name of each company whose employees earn a higher salary, on average, than the average salary at “First Bank Corporation”.

employee (ID, person_name, street, city)
works (ID, company_name, salary)
company (company_name, city)
manages (ID, manager_id)

Figure 3.19 Employee database.

- select E.ID, E.person_name, E.city*
from employee as E, works as W
where E.ID = W.ID and W.company_name = "First Bank Corporation"
- select E.ID, E.person_name, E.city*
from employee as E, works as W
where E.ID = W.ID
and W.company_name = "First Bank Corporation"
and W.salary > 10000
- select ID*
from works
where company_name < > "First Bank Corporation"
- select ID*
from works
where salary > all(select salary
from works
where company_name = "Small Bank Corporation"

第二次作业 3.9

- e. Assume that companies may be located in several cities. Find the name of each company that is located in every city in which “Small Bank Corporation” is located.
- f. Find the name of the company that has the most employees (or companies, in the case where there is a tie for the most).
- g. Find the name of each company whose employees earn a higher salary, on average, than the average salary at “First Bank Corporation”.

- e.

```
select C.company_name  
from company as C  
where not exists (select city  
                  from company  
                  where company_name = "Small Bank Corporation"  
                  except (select D.city from company as D  
                          where D.company_name = C.company_name))
```
- f.

```
select company_name  
from works  
group by company_name  
having count(distinct ID) >= all(  
                  select count(distinct ID) from works group by company_name)
```
- g.

```
select company_name  
from works  
group by company_name  
having avg(salary) > (select avg(salary) from works  
                  where company_name = "First Bank Corporation")
```

第二次作业 3.15

- 3.15** Consider the bank database of Figure 3.18, where the primary keys are underlined. Construct the following SQL queries for this relational database.
- Find each customer who has an account at *every* branch located in “Brooklyn”.
 - Find the total sum of all loan amounts in the bank.
 - Find the names of all branches that have assets greater than those of at least one branch located in “Brooklyn”.

branch(*branch_name*, *branch_city*, *assets*)
customer (*ID*, *customer_name*, *customer_street*, *customer_city*)
loan (*loan_number*, *branch_name*, *amount*)
borrower (*ID*, *loan_number*)
account (*account_number*, *branch_name*, *balance*)
depositor (*ID*, *account_number*)

Figure 3.18 Banking database.

第二次作业 3.15

- a. Find the ID of each customer of the bank who has an account but not a loan.
- b. Find the ID of each customer who lives on the same street and in the same city as customer '12345'.
- c. Find the name of each branch that has at least one customer who has an account in the bank and who lives in “Harrison”.

a. *select distinct C.customer_name*
from customer as C
where not exists (

(select branch_name
from branch

where branch_city = 'Brooklyn')

except

(select branch_name

from account join depositor using(account_number)

where C.ID = depositor.ID))

b. *select sum(amount)*
from loan

c. *select branch_name*
from branch

where assets > some

(select assets

from branch

where branch_city = “Brooklyn”)

第二次作业 3.16

3.16 Consider the employee database of Figure 3.19, where the primary keys are underlined. Give an expression in SQL for each of the following queries.

- Find ID and name of each employee who lives in the same city as the location of the company for which the employee works.
- Find ID and name of each employee who lives in the same city and on the same street as does her or his manager.
- Find ID and name of each employee who earns more than the average salary of all employees of her or his company.
- Find the company that has the smallest payroll.

employee (ID, person_name, street, city)
works (ID, company_name, salary)
company (company_name, city)
manages (ID, manager_id)

Figure 3.19 Employee database.

a. *select ID, person_name*
from (*employee natural join works*) *natural join*
company

c. *select ID, person_name*
from *works as A*
where salary > (select avg(salary)
from works as W
where W. company_name = A. company_name)

d. *select company_name*
from *works*
group by company_name
having sum(salary) <= all(select sum(salary)
from works
group by company_name)

b. *select E.ID, E.person_name*
from employee as E, employee as D, manages
where E.ID = manages.ID
and D.ID = manages.manager_id
and E.street = D.street
and E.city = D.city

第二次作业 3.17

3.17 Consider the employee database of Figure 3.19. Give an expression in SQL for each of the following queries.

- Give all employees of “First Bank Corporation” a 10 percent raise.
- Give all managers of “First Bank Corporation” a 10 percent raise.
- Delete all tuples in the *works* relation for employees of “Small Bank Corporation”.

employee (ID, person_name, street, city)
works (ID, company_name, salary)
company (company_name, city)
manages (ID, manager_id)

Figure 3.19 Employee database.

- update works*
set salary = salary * 1.1
where company_name = "First Bank Corporation"
- update works*
set salary = salary * 1.1
where works.company_name = "First Bank Corporation"
and ID in (*select* distinct manager_id
from *manages*)
- delete from works*
where company_name = "Small Bank Corporation"

第二次作业 3.21

3.21 Consider the library database of Figure 3.20. Write the following queries in SQL.

- Find the member number and name of each member who has borrowed at least one book published by “McGraw-Hill”.
- Find the member number and name of each member who has borrowed every book published by “McGraw-Hill”.
- For each publisher, find the member number and name of each member who has borrowed more than five books of that publisher.
- Find the average number of books borrowed per member. Take into account that if a member does not borrow any books, then that member does not appear in the *borrowed* relation at all, but that member still counts in the average.

member(*memb_no*, *name*)
book(*isbn*, *title*, *authors*, *publisher*)
borrowed(*memb_no*, *isbn*, *date*)

Figure 3.20 Library database.

a. *select M.memb_no, M.name*
from member M, book B, borrowed C
where C.memb_no = M.memb_no
and B.isbn = C.isbn
and B.publisher = "McGraw-Hill"

b. *select M.memb_no, M.name*
from member M
where not exist (select isbn from book
where B.publisher = "McGraw-Hill")
except (select isbn from borrowed E
where E.memb_no = M.memb_no)

c. *select M.memb_no, M.name, B.publisher*
from member M, book B, borrowed C
where C.memb_no = M.memb_no
and B.isbn = C.isbn
group by B.publisher, M.memb_no, M.name
having count() > 5*

d. *Select avg(borrowed_count) as avg_borrowed_count*
from (select ,count(b.isbn) as borrowed_count
from member as M
left join borrowed as B on M.memb_no = B.memb_no
group by M.memb_no
)as member_borrowed_count

第三次作业 4.7

employee (ID, person_name, street, city)
works (ID, company_name, salary)
company (company_name, city)
manages (ID, manager_id)

Figure 4.12 Employee database.

1. "employee"表, 包括ID、人名、街道和城市列。ID是主键。
2. "works"表包括ID、公司名和薪资列。ID是一个外键, 引用"employee"表, company_name是一个外键, 引用"company"表。这确保了员工和所在公司之间的引用完整性。
3. "company"表包含公司名和城市列。公司名是主键。
4. "manages"表用于跟踪管理关系。它包括ID和manager_id列, 都是外键, 引用"employee"表。

第三次作业 4.7

```
CREATE TABLE employee(  
    ID INT PRIMARY KEY,  
    person_name VARCHAR(128) NOT NULL,  
    street VARCHAR(128),  
    city VARCHAR(64)  
)  
  
CREATE TABLE works(  
    ID INT PRIMARY KEY,  
    company_name VARCHAR(128) NOT NULL,  
    salary INT NOT NULL,  
    FOREIGN KEY (company_name) REFERENCES company(company_name)  
)  
  
CREATE TABLE company(  
    company_name VARCHAR(128) PRIMARY KEY,  
    city VARCHAR(64)  
)  
  
CREATE TABLE manages(  
    ID INT PRIMARY KEY,  
    manager_id INT,  
    FOREIGN KEY (manager_id) REFERENCES employee(ID)  
)
```

第三次作业 4.16

- 4.16** Write an SQL query using the university schema to find the ID of each student who has never taken a course at the university. Do this using no subqueries and no set operations (use an outer join).

classroom(building, room_number, capacity)

department(dept_name, building, budget)

course(course_id, title, dept_name, credits)

instructor(ID, name, dept_name, salary)

section(course_id, sec_id, semester, year, building, room_number, time_slot_id)

teaches(ID, course_id, sec_id, semester, year)

student(ID, name, dept_name, tot_cred)

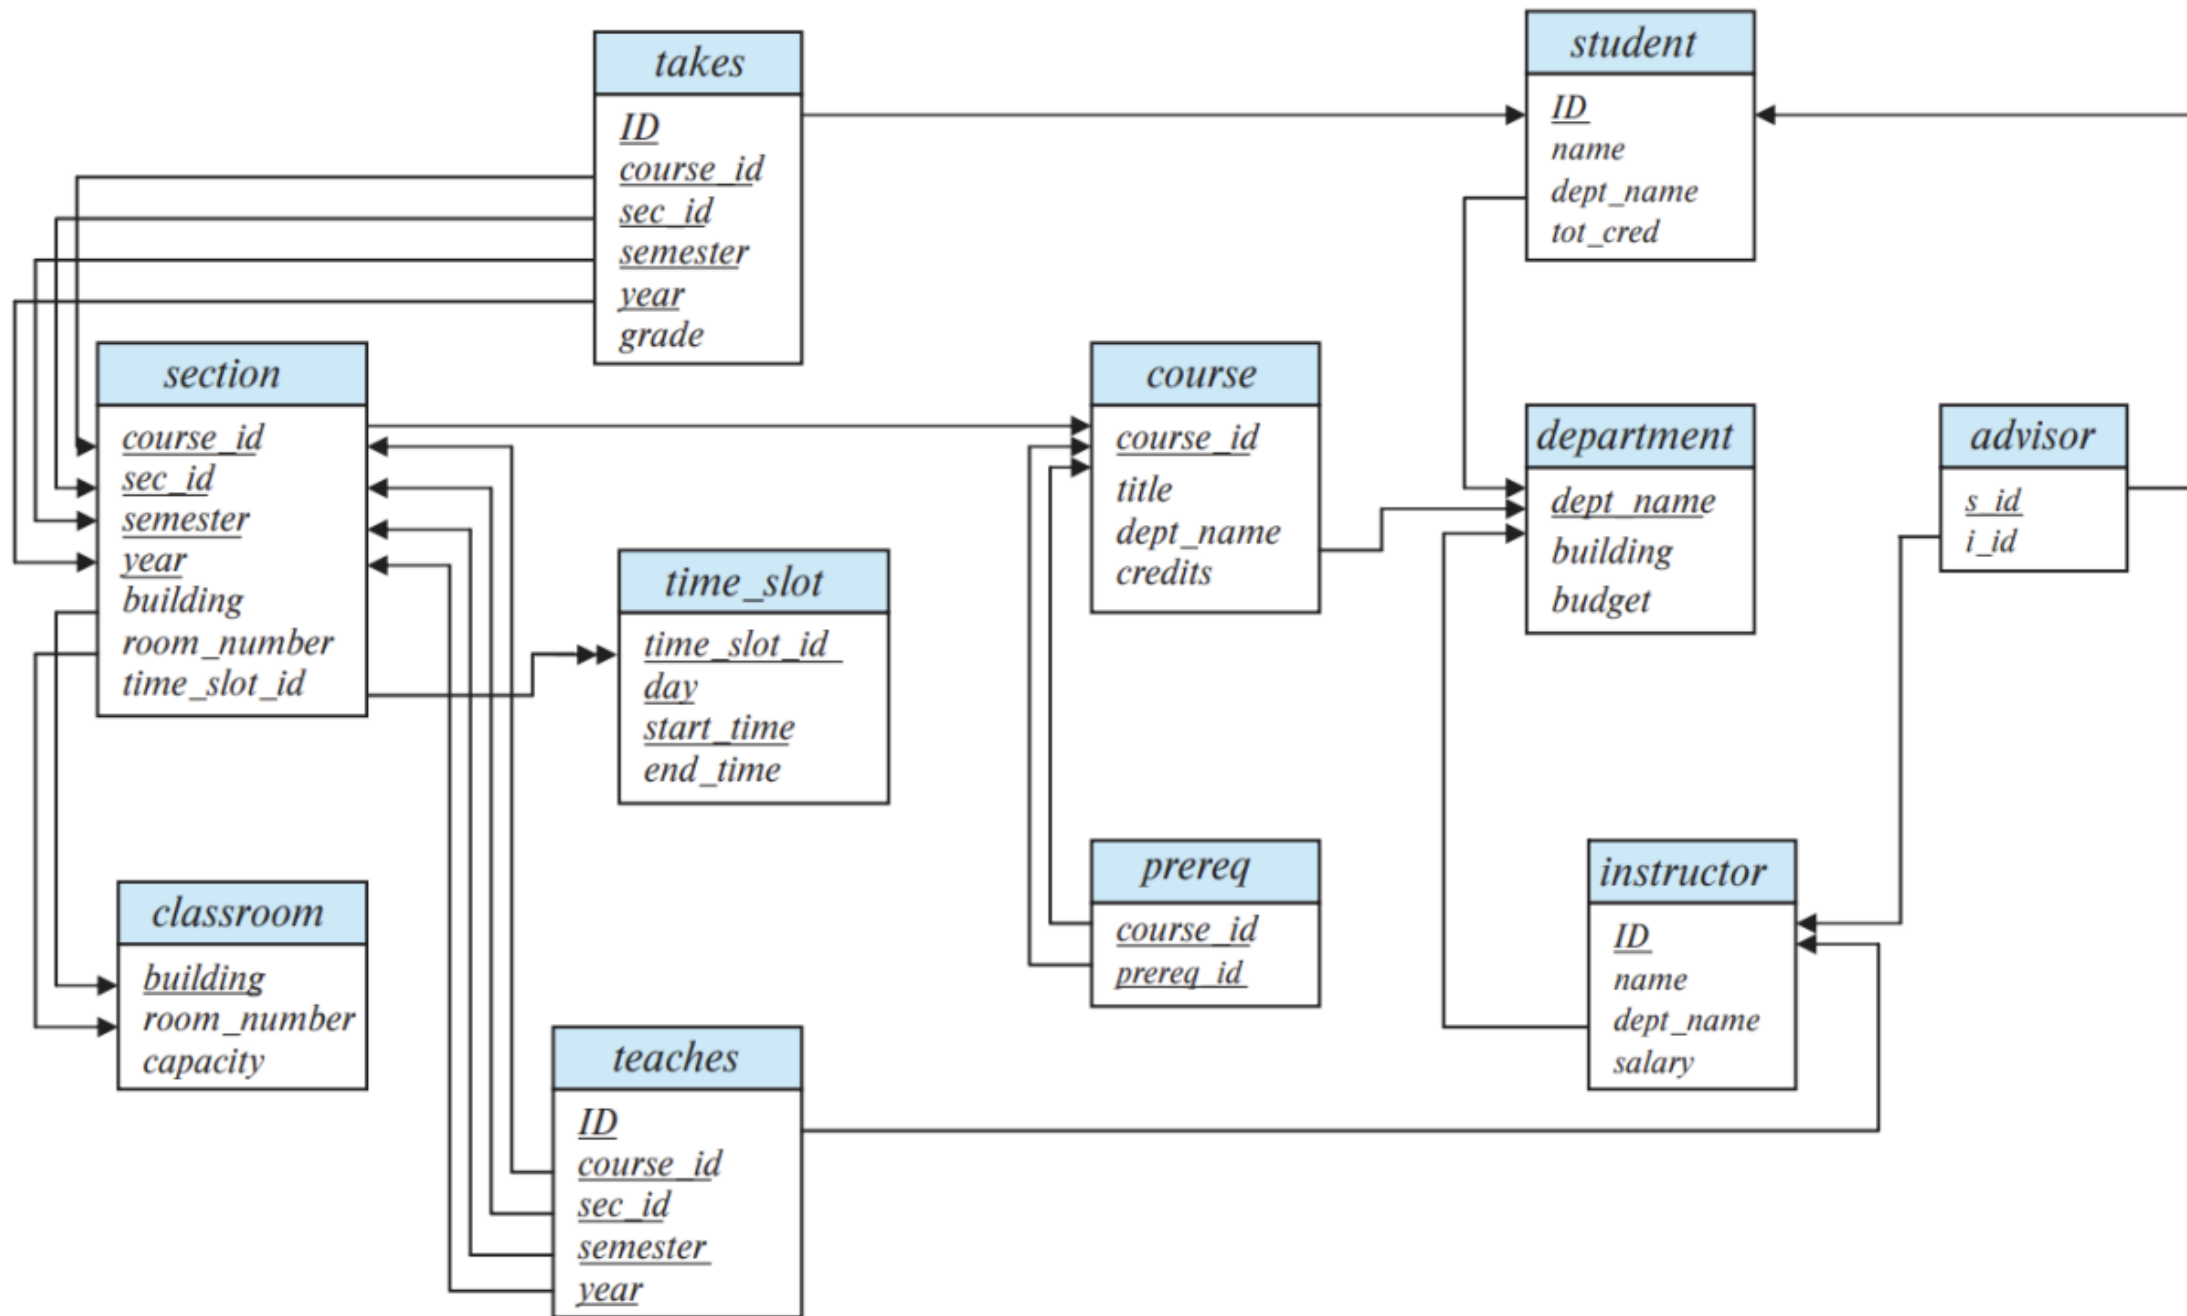
takes(ID, course_id, sec_id, semester, year, grade)

advisor(s_ID, i_ID)

time_slot(time_slot_id, day, start_time, end_time)

prereq(course_id, prereq_id)

第三次作业 4.16



```
SELECT s.ID
FROM student s LEFT OUTER JOIN takes t
ON s.ID = t.ID
WHERE t.ID IS NULL;
```

第四次作业 7.1

7.1 Suppose that we decompose the schema $R = (A, B, C, D, E)$ into

$$\begin{aligned} &(A, B, C) \\ &(A, D, E). \end{aligned}$$

Show that this decomposition is a lossless decomposition if the following set F of functional dependencies holds:

$$\begin{aligned} &A \rightarrow BC \\ &CD \rightarrow E \\ &B \rightarrow D \\ &E \rightarrow A \end{aligned}$$

根据无损分解的定义 $R_1 \cap R_2 \rightarrow R_1$ 或者 $R_1 \cap R_2 \rightarrow R_2$

因为 $R_1 = \{A, B, C\}$, $R_2 = \{A, D, E\}$, 所以 $R_1 \cap R_2 = \{A\}$

因为有 $A \rightarrow BC$, 所以可以得到 $A \rightarrow B$ 与 $A \rightarrow C$

因为 $A \rightarrow B$ 与 $B \rightarrow D$, 所以有 $A \rightarrow D$

因为 $A \rightarrow CD$ 与 $CD \rightarrow E$, 所以有 $A \rightarrow E$

综上所述 $A \rightarrow ABCDE$

结合 $R_1 \cap R_2 = \{A\}$

有 $R_1 \cap R_2 \rightarrow R_1$

所以这是一个无损分解

第四次作业 7.6

7.6 Compute the closure of the following set F of functional dependencies for relation schema $R = (A, B, C, D, E)$.

$$\begin{aligned}A &\rightarrow BC \\ CD &\rightarrow E \\ B &\rightarrow D \\ E &\rightarrow A\end{aligned}$$

List the candidate keys for R .

因为有 $A \rightarrow BC$, 所以可以得到 $A \rightarrow B$ 与 $A \rightarrow C$

因为 $A \rightarrow B$ 与 $B \rightarrow D$, 所以有 $A \rightarrow D$

因为 $A \rightarrow CD$ 与 $CD \rightarrow E$, 所以有 $A \rightarrow E$

综上所述 $A \rightarrow ABCDE$

因为 $E \rightarrow A$, 所以 $E \rightarrow ABCDE$

因为 $CD \rightarrow E$, 所以 $CD \rightarrow ABCDE$

因为 $B \rightarrow D$ 且 $BC \rightarrow CD$, 所以有 $BC \rightarrow ABCDE$

综上所述, 候选码有: A, E, CD, BC

第四次作业 7.27

7.27 Use Armstrong's axioms to prove the soundness of the decomposition rule.

- **Decomposition rule.** If $\alpha \rightarrow \beta\gamma$ holds, then $\alpha \rightarrow \beta$ holds and $\alpha \rightarrow \gamma$ holds.

of rules is called **Armstrong's axioms** in honor of the person who first proposed it.

- **Reflexivity rule.** If α is a set of attributes and $\beta \subseteq \alpha$, then $\alpha \rightarrow \beta$ holds.
- **Augmentation rule.** If $\alpha \rightarrow \beta$ holds and γ is a set of attributes, then $\gamma\alpha \rightarrow \gamma\beta$ holds.
- **Transitivity rule.** If $\alpha \rightarrow \beta$ holds and $\beta \rightarrow \gamma$ holds, then $\alpha \rightarrow \gamma$ holds.

因为 $\alpha \rightarrow \beta\gamma$

根据增强规则：有 $\alpha\beta \rightarrow \beta\gamma$

根据反射性规则：因为有 $\beta \subseteq \gamma\beta$, 所以 $\gamma\beta \rightarrow \beta$

根据传递性规则，因为 $\alpha \rightarrow \beta\gamma$ 且 $\gamma\beta \rightarrow \beta$, 所以有 $\alpha \rightarrow \beta$

证明 $\alpha \rightarrow \gamma$ 时同理，故不再赘述。

第四次作业 7.30

7.30 Consider the following set F of functional dependencies on the relation schema (A, B, C, D, E, G) :

$$\begin{aligned}A &\rightarrow BCD \\ BC &\rightarrow DE \\ B &\rightarrow D \\ D &\rightarrow A\end{aligned}$$

a. Compute B^+ .

因为 $B \rightarrow D$ 且 $D \rightarrow A$ 所以有 $B \rightarrow ABD$

因为 $A \rightarrow BCD$, 所以有 $B \rightarrow ABCD$

因为 $BC \rightarrow DE$, 所以有 $B \rightarrow ABCDE$

综上所述, $B^+ = \{A, B, C, D, E\}$

第四次作业 7.30

b. Prove (using Armstrong's axioms) that AG is a superkey.

$A \rightarrow BCD$, 由分解律: $A \rightarrow BC$

因为 $BC \rightarrow DE$, 由传递律: $A \rightarrow DE$

由合并律: $A \rightarrow BCDE$

由增广律: $AG \rightarrow ABCDEG$

所以 AG 是超码

第四次作业 7.30

c. Compute a canonical cover for this set of functional dependencies F ; give each step of your derivation with an explanation.

- 对于 $A \rightarrow BCD$, D 是无关属性, 因为在 $(F - \{A \rightarrow BCD\}) \cup (A \rightarrow BC)$ 中, $A \rightarrow BC$ 和 $B \rightarrow D$ 可以推出 $A \rightarrow D$
- 对于 $BC \rightarrow DE$, D 是无关属性, 因为在 $(F - \{BC \rightarrow DE\}) \cup (BC \rightarrow E)$ 中, $B \rightarrow D$ 可以推出 $BC \rightarrow D$
- 对于 $BC \rightarrow E$, C 是无关属性, 因为由 $B \rightarrow D$ 、 $D \rightarrow A$ 、 $A \rightarrow BC$ 可以推出 $B \rightarrow C$, 而由增广律可以得到 $B \rightarrow BC$, 再由合并律即可得到 $B \rightarrow E$, 因此 C 是无关属性
- 此时 F 可被简化为 $\{A \rightarrow BC, B \rightarrow E, B \rightarrow D, D \rightarrow A\}$
- 由于有两个左侧为 B 的函数依赖, 则将其合并, 可以得到 F 的正则覆盖为:

$$A \rightarrow BC$$

$$B \rightarrow DE$$

$$D \rightarrow A$$

第四次作业 7.30

- d. Give a 3NF decomposition of the given schema based on a canonical cover.

正则覆盖中没有多余的函数依赖

所以属性集是正则覆盖中其它函数依赖组成的子集

所以三个函数依赖都有自己的关系

$$R_1(A, B, C)$$

$$R_2(B, D, E)$$

$$R_3(D, A)$$

AG 是超码且上述关系没有原关系的超码，所以需要添加

$$R_4(A, G)$$

最终结果为

$$R_1(A, B, C)$$

$$R_1(B, D, E)$$

$$R_3(D, A)$$

$$R_4(A, G)$$

第四次作业 7.30

- Given the relational schema $R\langle U, F \rangle$, $U=\{A,B,C,D,E\}$, $F=\{AC \rightarrow BD, B \rightarrow C, C \rightarrow D, B \rightarrow E\}$
 - a) Use **Armstrong axioms** and related rules to prove the functional dependency $AC \rightarrow E$
 - b) Compute $(A)^+$ and $(AC)^+$
 - c) Find a **canonical cover** F_c of F
 - d) Find **all candidate keys**, and point out R is in **which normal form**
 - e) **Decompose R into 3NF**, which the decomposition is **lossless-join** and **dependency preserving**.
 - f) Give related explanation or proof that the above decomposition is **lossless-join** and **dependency preserving**
 - g) *Decompose the relation into relations in **BCNF**

a) $\therefore AC \rightarrow BD$
 $\therefore AC \rightarrow B$
 $\therefore \text{条件 } B \rightarrow E$
 $\therefore AC \rightarrow E$

b) $(A)^+ = \{A\}$
 $\therefore AC \rightarrow BD \quad B \rightarrow E \therefore AC \rightarrow E$
 $\therefore (AC)^+ = \{A, B, C, D, E\}$

第四次作业 7.30

c) 将 $AC \rightarrow BD$ 分解为 $AC \rightarrow B$ 和 $AC \rightarrow D$
得到: $AC \rightarrow B$ $AC \rightarrow D$ $B \rightarrow C$ $C \rightarrow D$ $B \rightarrow E$
由 $AC \rightarrow D$ 和 $C \rightarrow D$ 可以去掉 A : $C \rightarrow D$
因此正则覆盖为:
 $\{AC \rightarrow B, B \rightarrow CE, C \rightarrow D\}$

d) $A^+ = \{A\}$
 $B^+ = \{B, C, D, E\}$ $AB^+ = \{A, B, C, D, E\}$
 $C^+ = \{C, D\}$ $AC^+ = \{A, C, B, D, E\}$
 $D^+ = \{D\}$
 $E^+ = \{E, A\}$ 满足 $1F$ 范式

e) $R_1 = (A, B, C)$
 $R_2 = (C, D)$
 $R_3 = (B, C, E)$

第四次作业 7.30

f)

	A	B	C	D	E
ABC	a_1	a_2	a_3	b_{14} ^{a_4}	b_{15} ^{a_5}
CD	b_{21}	b_{22}	a_3	a_4	b_{25}
BCE	b_{31}	a_2	a_3	b_{34} ^{a_4}	a_5

$$B \rightarrow E: b_{15} \rightarrow a_5$$

$$C \rightarrow D: b_{14} \rightarrow a_4 \quad b_{34} \rightarrow a_4$$

e)

分解结果为 $\{A, B, C\}, \{B, C\}, \{B, E\}, \{C, D\}$

第五次作业

- Q1: Construct a **B⁺-tree** from an empty tree. Each node can hold **four pointers**
 - The sequential values to be inserted are: 10, 7, 12, 5, 9, 15, 30, 23, 17, 26
 - Then delete 9, 10, 15, respectively
 - Please give the B⁺ trees after each insertion and each deletion

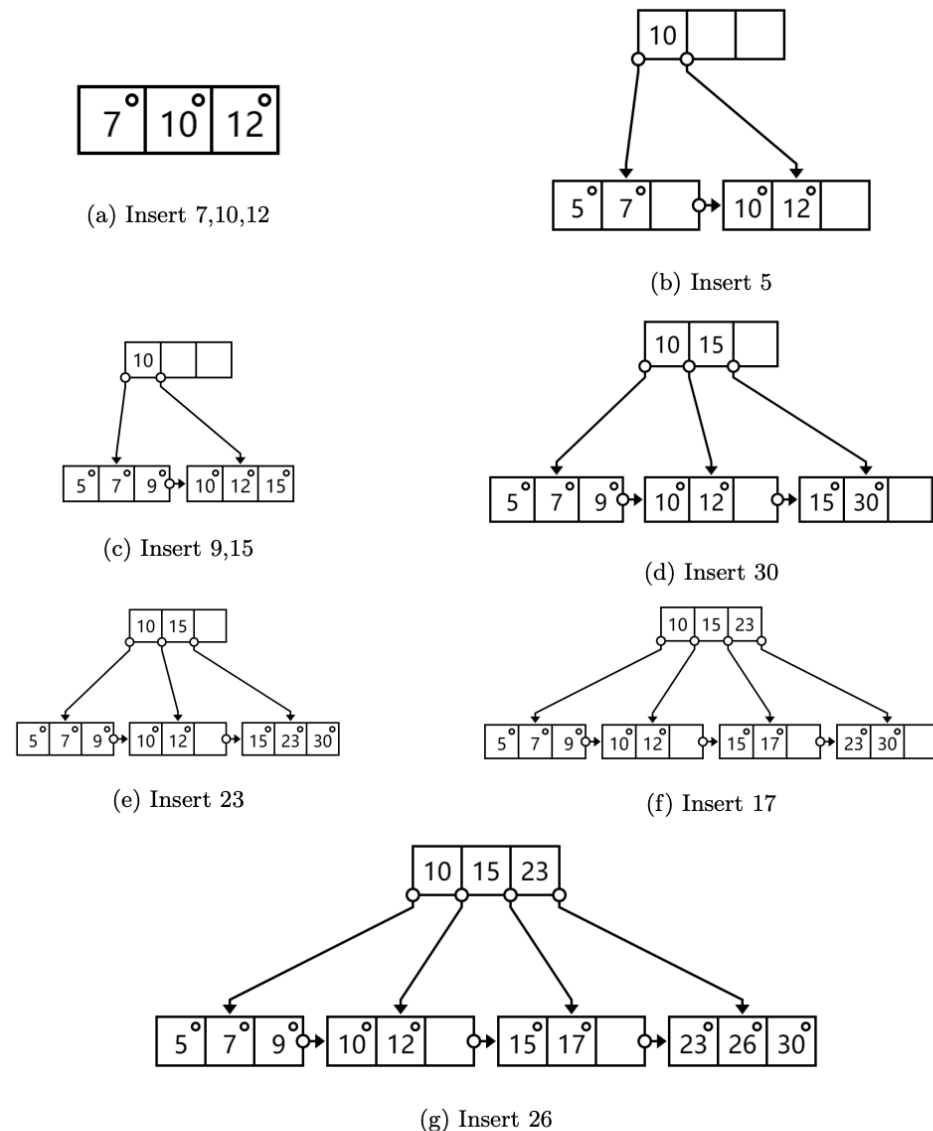


Figure 1.1.1: B⁺ tree after sequential insertion (Part I)

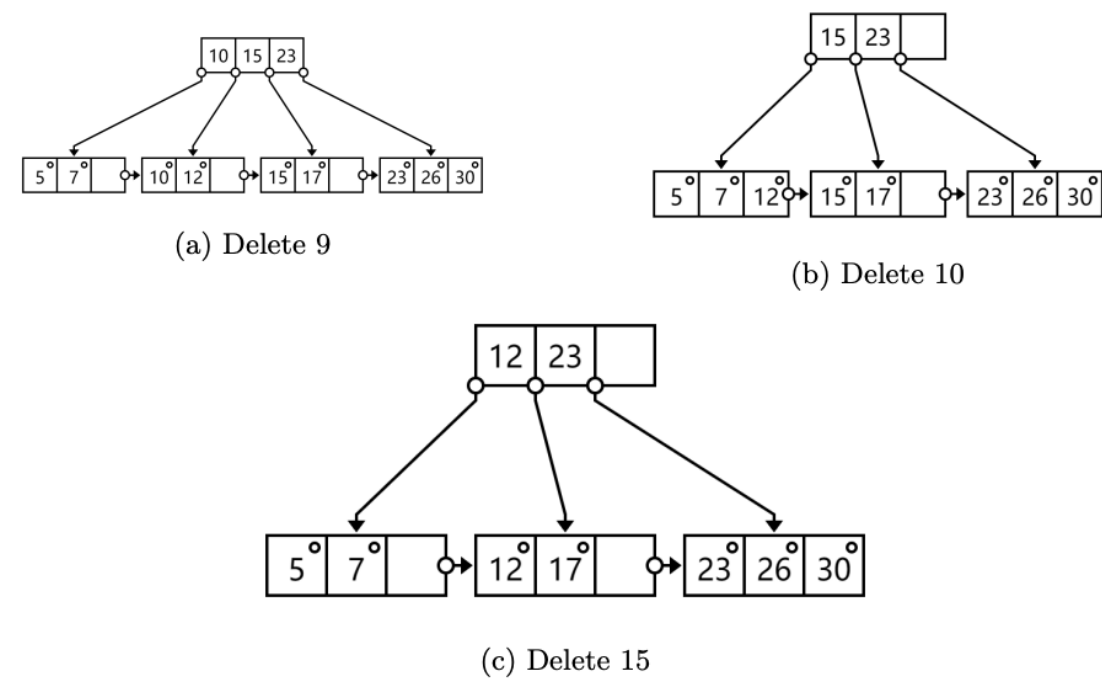


Figure 1.1.2: B⁺ tree after sequential deletion (Part I)

第五次作业

- Q2: Compare **B⁺-tree** and **B-tree** and describe their difference

1. 节点存储数据的方式:

B-树: 内部节点和叶子节点都可以存储数据。每个节点中的关键字信息包含关键码和指向子节点的指针, 这些关键码和指针用于索引和查找数据。

B⁺树: 内部节点不存储数据, 只作为索引存在。所有数据都存储在叶子节点中, 且叶子节点之间通过链表相连, 这使得范围查询等操作更为高效。

2. 查找过程:

B-树: 在查找过程中, 如果找到了具体的数值, 就会结束查找。

B⁺树: 由于所有数据都存储在叶子节点中, 所以查找过程会一直进行到叶子节点, 通过索引找到叶子节点中的数据才结束。

3. 关键字出现频率:

B-树: 中任何一个关键字出现且只出现在一个节点中。

B⁺树: 由于叶子节点之间存在链表连接, 所以关键字可能在多个叶子节点中出现。

4. 磁盘 I/O 与查询时间复杂度:

B-树: 查询时间复杂度不固定, 与 Key 在树中的位置有关。在某些情况下, 如果 Key 恰好位于根节点或接近根节点的位置, 查询时间可能会非常快 ($O(1)$)。

B⁺树: 由于叶子节点的数据都是使用链表连接起来的, 并且在磁盘中是顺序存储的, 所以当读到某个值时, 磁盘预读原理会提前把这些数据都读进内存。这使得范围查询和排序等操作非常高效。同时, B⁺树的查询时间复杂度是固定的, 为 $O(\log n)$ 。



5. 应用场景:

B-树: 由于内部节点也可以存储数据, 所以 B-树在需要频繁进行插入、删除等操作的场景下表现较好。它也被广泛应用于文件系统和数据库索引中。


B⁺树: 由于所有数据都存储在叶子节点中, 并且叶子节点之间通过链表相连, 这使得 B⁺树在范围查询和排序等操作中表现出色。因此, B⁺树更适合用于需要频繁进行范围查询和排序的场景, 如数据库索引等。

第五次作业

- For the B+ tree index, please give the B+ trees after insertion of search key 6 and 8:




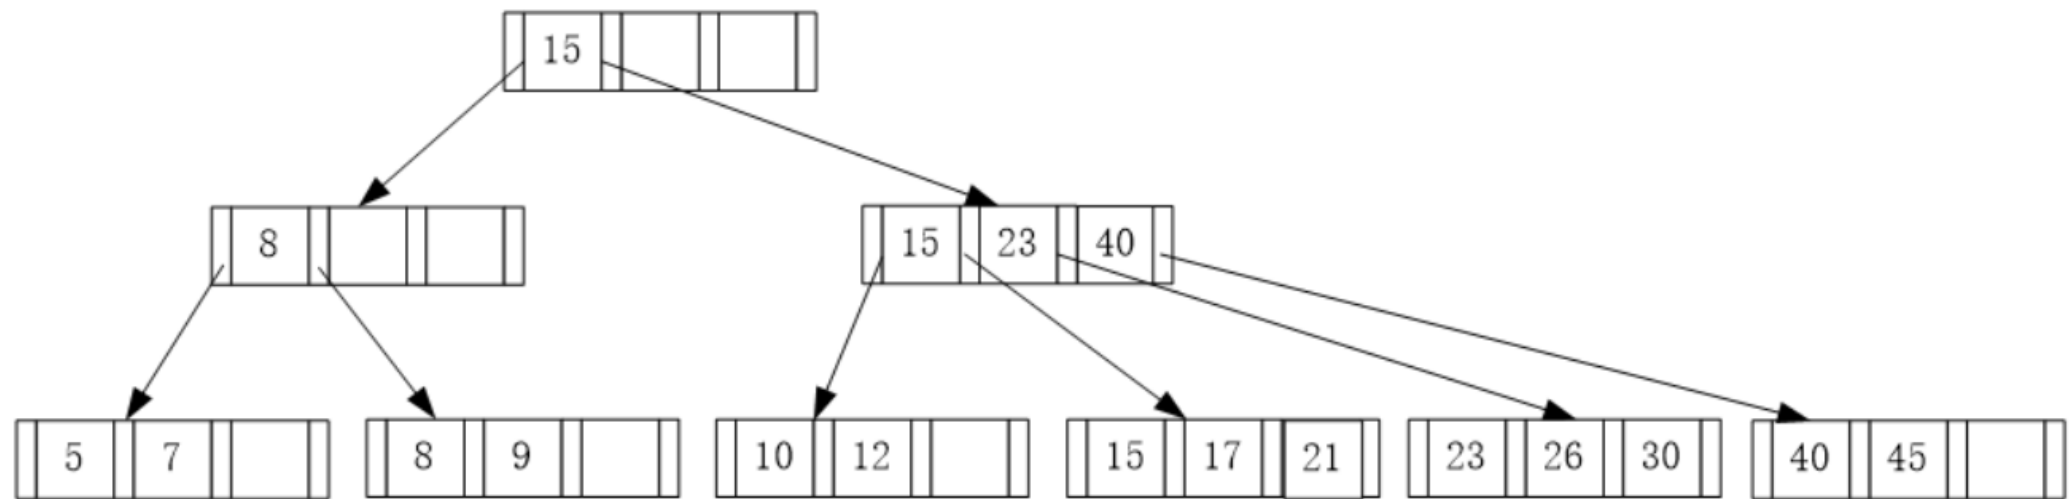
(a) Insert 6




(b) Insert 8

第五次作业

- Pease give the B+ trees after deletion of search keys both 8 and 9:



(a) Delete 8



(b) Delete 9

第六次作业

16.16 Suppose that a B⁺-tree index on (*dept_name*, *building*) is available on relation *department*. What would be the best way to handle the following selection?

$$\sigma_{(building < \text{"Watson"}) \wedge (budget < 55000) \wedge (dept_name = \text{"Music"})}(department)$$

branch(*branch_name*, *branch_city*, *assets*)
customer (*customer_name*, *customer_street*, *customer_city*)
loan (*loan_number*, *branch_name*, *amount*)
borrower (*customer_name*, *loan_number*)
account (*account_number*, *branch_name*, *balance*)
depositor (*customer_name*, *account_number*)

Figure 16.9 Banking database.

第六次作业

在给定的条件下，最好的处理方式是利用 B+-树索引来优化查询。这是因为 B+-树索引在处理范围查询和精确匹配查询时非常高效。

假设有一个 B+-树索引存在于 (dept name, building) 上，可以采取以下步骤来处理查询 σ (building < "Watson") \wedge (budget < 55000) \wedge (dept name = "Music")(department):

1. 索引扫描：首先利用 dept name 列来进行精确匹配查询，因为 B+-树索引的首要列是 dept name。查询 dept name = "Music"，找到所有属于 Music 系的记录。
2. 范围查询：对于第一步中匹配的记录，使用 building < "Watson" 条件进行范围查询。由于 building 是索引的第二个列，可以在索引上高效地执行这个范围查询。
3. 过滤其他条件：对满足前两个条件的记录，在主表中提取出来，然后再检查 budget < 55000 的条件。这一步可能需要访问主表，但由于前两步已经大大减少了需要检查的记录数，这个操作会相对高效。